It's been a while since I took part in a Tidy Tuesday, but having not played with any R code for a while I got the urge to create some utterly pointless plots this weekend, so here we are. This week's dataset is looking at world records in Mario Kart 64, which despite being over two decades old (don't let it set in), is still active in the videogame speedrunning community.

First up, a ridgeline plot of the density of world record activity per track over the years:

I'd seen ridgeplots quite a lot when browsing #TidyTuesday on Twitter so I figured it was about time I gave one a shot myself. I used the ggridges package and it was fun to finally have an excuse to try it out! Looking at the plot, it seems like there was a resurgence in world record activity later in the games life – most likely due to more widespread internet access and the growth of the speedrunning scene in general as a result, I imagine.

Second, a plot showing the total amount of time it would take you to complete all the tracks in Mario Kart 64 at the fastest world record pace each year:

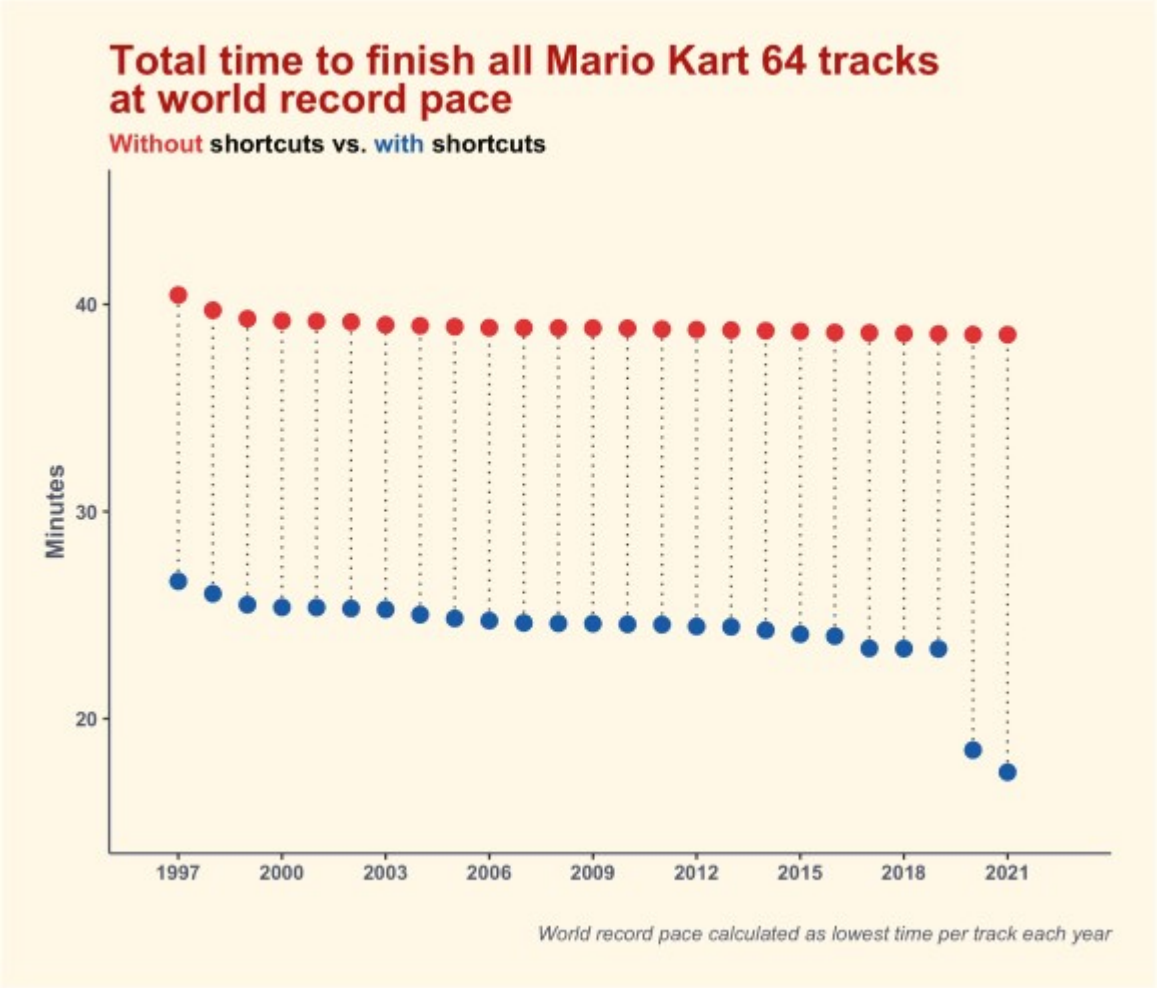
Total time to finish all Mario Kart 64 tracks at world record pace
Without shortcuts vs. with shortcuts
Minutes
40
30
20
1997
2000
2003
2006
2009
2012
2015
2018
2021
World record pace calculated as lowest time per track each year